

3 Repeated Measures Analysis of Variance

3.1 Introduction

As we have discussed, many approaches have been taken in the literature to specifying **statistical models** for longitudinal data. Within the framework of a specific model, the questions of scientific interest are interpreted and represented formally, and associated with the model are **statistical methods** that allow the questions to be addressed. Different models embodying different assumptions and taking different perspectives (e.g., SS vs. PA) lead to possibly different characterizations of the questions and different methods.

We begin our review of approaches by considering two statistical models that form the basis for what we have called **classical methods**. These methods are appropriate for **continuous outcome** or outcomes that can be viewed as approximately continuous and that are reasonably thought to be approximately **normally distributed**.

The models have several limitations relative to those underlying the more **modern** approaches that we discuss in subsequent chapters.

- The models are really only applicable in the case of **balanced** data; that is, where the responses on each individual are recorded at the **same time points**, with no departures from these times or missing data. Thus, in this chapter, we assume that each individual is observed at the same n time points t_1, \dots, t_j , say, and each has an associated n -dimensional response vector, where the j th element corresponds to the response at time t_j .
- The models adopt a representation of the **overall population mean** of a response vector that is **simplistic** in the sense that it does not recognize the fact that the mean response might exhibit a **systematic trajectory** over **continuous** time, such as a **straight line**. In particular, the mean ordinarily is represented using the notation that is commonplace when developing **analysis of variance** methods. Time and **among-individual covariates** are viewed as **categorical factors** with a small number of **levels**. The treatment of time as a categorical factor is particularly restrictive.

- The models do not accommodate straightforwardly incorporation of **covariate information** beyond time and among-individual categorical factors. In our discussion, we restrict attention to a **single** among-individual factor such as **group membership**; e.g., gender in the dental study or dose in the guinea pig diet study in **EXAMPLE 2** of Section 1.2. Although the models can be generalized to more than one such factor (e.g., genotype and weather pattern in the soybean growth study of **EXAMPLE 3** of Section 1.2), we do not consider this, deferring to the more modern approaches we discuss later that allow much greater flexibility.
- The associated analysis methods are focused almost entirely on **hypothesis testing**. Thus, they do not readily accommodate questions regarding the nature of features of mean trajectories; e.g., in the dental study, the values of the **slopes** characterizing **rate of change** of assumed **straight line** population mean trajectories for boys and girls. That is, **estimation** of quantities like these is not straightforward within these modeling frameworks.

As we demonstrate, the models also embody assumptions on the **overall covariance structure** of a response vector that are possibly **too restrictive** or **too general**.

- The statistical model underlying **univariate repeated measures analysis of variance** (ANOVA) methods is derived from a **SS** perspective. As we noted in Section 2.5, it induces a model for the overall covariance pattern that has a **compound symmetric** correlation structure, which may or may not be a plausible model.
- The statistical model underlying **multivariate repeated measures analysis of variance** methods arises from **PA** perspective. **No specific systematic assumption** is made on the overall covariance pattern, so that it is regarded as **completely unstructured**. If correlation **does** exhibit a simpler pattern, these methods could be **inefficient**.

We first review the **univariate** methods, followed by the **multivariate** approach. Our discussion is limited to the basic elements and thus is not meant to be **comprehensive**. Rather, it is meant only to provide an appreciation of why statistical practice has moved toward favoring the modern methods discussed in subsequent chapters. Accordingly, we simply present results and do not offer detailed derivations or proofs.

3.2 Univariate repeated measures analysis of variance

We first discuss the basic model underlying univariate repeated measures ANOVA methods in the *classical notation*. In particular, we present the usual “**one way**” model, where there is a **single among-individual factor** such as gender in the dental study of **EXAMPLE 1** of Section 1.2 or vitamin E dose level in the guinea pig growth study of **EXAMPLE 2**.

BASIC SET-UP:

- The model assumes that the data arise from a study in which individuals are **randomized** or naturally belong to one of $g \geq 1$ groups; the group variable is often referred to as the **between- or among-units** factor. Thus, in the dental study, $g = 2$ genders; in the guinea pig growth study, $g = 3$ dose groups.

From the point of view of the general notation introduced in Section 2.2, the model thus accommodates a **single** scalar, categorical **among-individual covariate** with g possible values. The model does not allow for **within-individual covariates**.

- The response is recorded on each of n occasions or under each of n conditions. In a longitudinal study, this is usually “**time**” but could be another repeated measurement condition. E.g., if men are randomized into two groups, regular and modified diet, the response might be maximum heart rate after separate occasions during which each spent 10, 20, 30, 45, and 60 minutes walking on a treadmill. We use the generic term **time**; this factor is often referred to in the classical literature as the **within-units** factor. In the dental study, this is age ($n = 4$); in the guinea pig study, weeks ($n = 6$).

Thus, from the point of view of the **conceptual framework** discussed in Section 2.3, the model does not acknowledge explicitly that there is an underlying process in **continuous time** and that there could be values of the response at times **other** than these n occasions.

- As noted above, we consider only the case where there is a **single group factor**. However, it is straightforward to extend the development to the case where the groups are determined by a **factorial design**; e.g. if in the guinea pig study there had been $g = 6$ groups, determined by the factorial arrangement of 3 doses and 2 genders.

NOTATION AND MODEL: We present the model first using the *classical notation* and then demonstrate how it can be expressed in the notation introduced in Chapter 2. Define

$Y_{h\ell j}$ = response on individual h in the ℓ th group at time j .

- $h = 1, \dots, r_\ell$, where r_ℓ denotes the number of units in group ℓ . Thus, in this notation, h indexes units **within** a particular group; $\ell = 1, \dots, g$ indexes groups; and $j = 1, \dots, n$ indexes the levels of time. Note then that a specific individual is **uniquely** identified by the indices (h, ℓ) .
- The total number of individuals is $m = \sum_{\ell=1}^g r_\ell$. Each is observed at n time points.

The *classical model* for $Y_{h\ell j}$ is then given by

$$Y_{h\ell j} = \mu + \tau_\ell + b_{h\ell} + \gamma_j + (\tau\gamma)_{\ell j} + e_{h\ell j}. \quad (3.1)$$

In the usual terminology accompanying classical ANOVA methods

- μ is an “overall mean,” τ_ℓ is the **fixed deviation** from the overall mean associated with being in group ℓ , γ_j is the **fixed deviation** associated with time j , and $(\tau\gamma)_{\ell j}$ is an **additional fixed deviation** associated with group ℓ and time j ; that is, $(\tau\gamma)_{\ell j}$ is the **interaction** effect for group ℓ and time j .
- Thus, as we demonstrate explicitly below, the **overall population mean response** for the ℓ th group at time j is represented as

$$\mu + \tau_\ell + \gamma_j + (\tau\gamma)_{\ell j}.$$

- $b_{h\ell}$ is a **random effect** assumed to be **independent** of the among-individual covariate group with conditional (on group) mean equal to the unconditional mean $E(b_{h\ell}) = 0$ characterizing how the “**inherent mean**” for the h th individual in group ℓ **deviates** from the **overall population mean**. Thus, (3.1) represents the **inherent (conditional) mean** for individual (h, ℓ) as

$$\mu + \tau_\ell + \gamma_j + (\tau\gamma)_{\ell j} + b_{h\ell}, \quad (3.2)$$

and $b_{h\ell}$ characterizes **among-individual** behavior.

- $e_{h\ell j}$ is a **within-individual** deviation representing the net effect of realizations and measurement error, **independent** of the among-individual covariate group, with conditional (on group) mean equal to the unconditional mean $E(e_{h\ell j}) = 0$. This is often called the “**random error**,” but as we have remarked previously we prefer the term **within-individual deviation** to reflect the fact that it embodies more than just measurement error.

Some observations are immediate.

- Model (3.1) has the **same form** as the statistical model for observations arising from an experiment conducted according to a **split plot** design. Thus, as we show shortly, the analysis is **identical** to that of a split plot experiment; however, the **interpretation** and further analyses are different.
- The **actual values** of the times (e.g. ages 8, 10, 12, 14 in the dental study) **do not** appear explicitly in the model. Rather, a separate deviation parameter γ_j and interaction parameter $(\tau\gamma)_{\ell j}$ is associated with each time. Thus, the model takes no account of where the times of observation are temporally; e.g. are they **equally-spaced** ?

Because (3.1) is a **linear model**, as discussed in Section 2.4, we can view it as a SS or PA model.

- From a SS perspective, as in (3.2),

$$\mu + \tau_{\ell} + \gamma_j + (\tau\gamma)_{\ell j} + b_{h\ell}$$

represents the inherent mean trend for the h th individual in group ℓ at time j . Note this assumes that the inherent mean for a given individual (h, ℓ) deviates from the overall population mean by the **same amount**, $b_{h\ell}$, at each time j . Thus, this model implies that if an individual is “high” relative to the overall mean response at time j , the individual is “high” at all other times.

This is often not a reasonable assumption. For example, consider the the conceptual representation in Figure 2.2. This assumption might be reasonable for the two uppermost individuals in panel (b), as the “inherent trends” for these are roughly parallel to the overall mean response trajectory. However, it is clearly not appropriate for the lowermost unit.

- Taking a PA perspective, write (3.1) as

$$Y_{h\ell j} = \underbrace{\mu + \tau_{\ell} + \gamma_j + (\tau\gamma)_{\ell j}}_{\mu_{\ell j}} + \underbrace{b_{h\ell} + e_{h\ell j}}_{\epsilon_{h\ell j}}. \quad (3.3)$$

In (3.3), $\epsilon_{h\ell j} = b_{h\ell} + e_{h\ell j}$ is the overall deviation reflecting aggregate deviation from this mean due to **among-** and **within-individual** sources.

Because $b_{h\ell}$ and $e_{h\ell j}$ have mean 0 (conditional on the among-individual covariate group and unconditionally), it follows that

$$E(Y_{h\ell j}) = \mu_{\ell j} = \mu + \tau_{\ell} + \gamma_j + (\tau\gamma)_{\ell j},$$

the overall population mean for the ℓ th group at the j th time.

CONVENTION: Henceforth, we write $\mathcal{N}(\mu, \sigma^2)$ to denote a **univariate normal distribution** with mean μ and variance σ^2 . We write $\mathcal{N}(\mu, \mathbf{V})$ to denote a **multivariate normal distribution** with mean vector μ and covariance matrix \mathbf{V} . The meaning (univariate or multivariate) is ordinarily clear from the context.

NORMALITY AND INDEPENDENCE ASSUMPTIONS: The model is completed by standard assumptions on the random deviations $b_{h\ell}$ and $e_{h\ell j}$, which lead to an assumption on the form of the overall pattern of variance and correlation.

- $b_{h\ell} \sim \mathcal{N}(0, \sigma_b^2)$ and are **independent** for all h and ℓ , so that where any individual “sits” in the population is unrelated to where others “sit.” The fact that this normal distribution is **identical** for $\ell = 1, \dots, g$ reflects the assumption that $b_{h\ell}$ is independent of the among-individual covariate group, so that **among-individual variation** is **the same** in all g populations. The **variance component** σ_b^2 represents the common magnitude of among-individual variation.
- $e_{h\ell j} \sim \mathcal{N}(0, \sigma_e^2)$ and are **independent** for all h, ℓ , and j . As for $b_{h\ell}$, that this normal distribution is the same for $\ell = 1, \dots, 1$ follows from the assumption that the $e_{h\ell j}$ are independent of the among-individual covariate group. Moreover, it also reflects the assumption that **within-individual variation** is **the same** at all observation times. Independence across j also implies that **within-individual correlation** across the observation times is **negligible**. The variance component σ_e^2 represents the magnitude of within-individual variation aggregated from all within-individual sources, namely, the realization process and measurement error.
- The $b_{h\ell}$ and $e_{h\ell j}$ are assumed to be **mutually independent** for all h, ℓ , and j . From the conceptual representation point of view, this says that deviations due to within-individual sources are of similar magnitude regardless of the magnitudes of the deviations $b_{h\ell}$ associated with the units on which the observations are made. This is often reasonable; however, as we will see later in the course, there are situations where it may not be reasonable.

VECTOR REPRESENTATION AND OVERALL COVARIANCE MATRIX: We can summarize the model for the responses for individual (h, ℓ) in the $(n \times 1)$ **random vector**

$$\begin{pmatrix} Y_{h\ell 1} \\ Y_{h\ell 2} \\ \vdots \\ Y_{h\ell n} \end{pmatrix} = \begin{pmatrix} \mu + \tau_\ell + \gamma_1 + (\tau\gamma)_{\ell 1} \\ \mu + \tau_\ell + \gamma_2 + (\tau\gamma)_{\ell 2} \\ \vdots \\ \mu + \tau_\ell + \gamma_n + (\tau\gamma)_{\ell n} \end{pmatrix} + \begin{pmatrix} b_{h\ell} \\ b_{h\ell} \\ \vdots \\ b_{h\ell} \end{pmatrix} + \begin{pmatrix} e_{h\ell 1} \\ e_{h\ell 2} \\ \vdots \\ e_{h\ell n} \end{pmatrix} = \begin{pmatrix} \mu_{\ell 1} \\ \mu_{\ell 2} \\ \vdots \\ \mu_{\ell n} \end{pmatrix} + \begin{pmatrix} \epsilon_{h\ell 1} \\ \epsilon_{h\ell 2} \\ \vdots \\ \epsilon_{h\ell n} \end{pmatrix}. \quad (3.4)$$

With $\mathbf{1}$ a $(n \times 1)$ vector of 1s, we write (3.4) compactly as

$$\mathbf{Y}_{h\ell} = \boldsymbol{\mu}_\ell + \mathbf{1}b_{h\ell} + \mathbf{e}_{h\ell} = \boldsymbol{\mu}_\ell + \boldsymbol{\epsilon}_{h\ell}. \quad (3.5)$$

Under the foregoing assumptions, it is clear that each $Y_{h\ell j}$ is **normally distributed** with

$$E(Y_{h\ell j}) = \mu_{\ell j} = \mu + \tau_\ell + \gamma_j + (\tau\gamma)_{\ell j}, \quad \text{so that} \quad E(\mathbf{Y}_{h\ell}) = \boldsymbol{\mu}_\ell,$$

$$\text{var}(Y_{h\ell j}) = \text{var}(b_{h\ell}) + \text{var}(e_{h\ell j}) + 2\text{cov}(b_{h\ell}, e_{h\ell j}) = \sigma_b^2 + \sigma_e^2.$$

(conditionally and unconditionally). Moreover, it is straightforward to show (try it) that

$$\text{cov}(Y_{h\ell j}, Y_{h'\ell'j'}) = \text{cov}(\epsilon_{h\ell j}, \epsilon_{h'\ell'j'}) = 0, \quad h \neq h',$$

where $\ell \neq \ell'$ or $\ell = \ell'$ and $j \neq j'$ or $j = j'$; i.e., the covariance between observations from two different units from the same or different groups at the same or different times is zero, which implies under normality that $Y_{h\ell j}$ and $Y_{h'\ell'j'}$ are **independent**.

Thus, under the assumptions of the model, for $\ell \neq \ell'$ or $\ell = \ell'$, the random vectors $\mathbf{Y}_{h\ell}$ and $\mathbf{Y}_{h'\ell'}$ are **independent**, showing that the model **automatically induces** the usual assumption that data vectors from different individuals are independent.

It is also straightforward to derive that

$$\begin{aligned} \text{cov}(Y_{h\ell j}, Y_{h\ell'j'}) &= \text{cov}(\epsilon_{h\ell j}, \epsilon_{h\ell'j'}) = E\{(Y_{h\ell j} - \mu_{\ell j})(Y_{h\ell'j'} - \mu_{\ell'j'})\} = E\{(b_{h\ell} + e_{h\ell j})(b_{h\ell'} + e_{h\ell'j'})\} \\ &= E(b_{h\ell}b_{h\ell'}) + E(e_{h\ell j}b_{h\ell'}) + E(b_{h\ell}e_{h\ell'j'}) + E(e_{h\ell j}e_{h\ell'j'}) = \sigma_b^2. \end{aligned}$$

Summarizing, we have that the m data vectors $\mathbf{Y}_{h\ell}$, $h = 1, \dots, r_\ell$, $\ell = 1, \dots, g$ are all independent and multivariate normal; that is, $\mathbf{Y}_{h\ell} \sim \mathcal{N}_n(\boldsymbol{\mu}_\ell, \mathbf{V})$, where

$$\mathbf{V} = \text{var}(\mathbf{Y}_{h\ell}) = \text{var}(\boldsymbol{\epsilon}_{h\ell}) = \begin{pmatrix} \sigma_b^2 + \sigma_e^2 & \sigma_b^2 & \cdots & \sigma_b^2 \\ \sigma_b^2 & \sigma_b^2 + \sigma_e^2 & \cdots & \sigma_b^2 \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_b^2 & \sigma_b^2 & \cdots & \sigma_b^2 + \sigma_e^2 \end{pmatrix}. \quad (3.6)$$

- This result follows directly from (3.5). Using the independence of $b_{h\ell}$ and $\mathbf{e}_{h\ell}$,

$$\text{var}(\mathbf{Y}_{h\ell}) = \text{var}(\epsilon_{h\ell}) = \text{var}(\mathbf{1}b_{h\ell}) + \text{var}(\mathbf{e}_{h\ell}) = \text{var}(b_{h\ell})\mathbf{1}\mathbf{1}' + \text{var}(\mathbf{e}_{h\ell}),$$

which, writing

$$\mathbf{1}\mathbf{1}' = \mathbf{J}_n = \begin{pmatrix} 1 & \cdots & 1 \\ 1 & \cdots & 1 \\ \vdots & \vdots & \vdots \\ 1 & \cdots & 1 \end{pmatrix} \text{ and } \text{var}(\mathbf{e}_{h\ell}) = \sigma_e^2 \mathbf{I}_n,$$

yields

$$\text{var}(\mathbf{Y}_{h\ell}) = \text{var}(\epsilon_{h\ell}) = \sigma_b^2 \mathbf{J}_n + \sigma_e^2 \mathbf{I}_n = \mathbf{V}, \quad (3.7)$$

where (3.7) is a compact expression for (3.6).

- From (3.6),

$$\text{corr}(Y_{h\ell j}, Y_{h\ell j'}) = \frac{\sigma_b^2}{\sigma_b^2 + \sigma_e^2}, \quad j \neq j'. \quad (3.8)$$

Thus, the overall correlation between any two observations in $\mathbf{Y}_{h\ell}$ is **the same** and equal to (3.8). The quantity (3.8) is called the **intraclass correlation** in some contexts.

- These results show that the model assumes that the **overall aggregate pattern of correlation** is **compound symmetric** or **exchangeable**. Note that, automatically, the correlation between any two observations in $\mathbf{Y}_{h\ell}$ is assumed to be **positive**, as ordinarily $\sigma_b^2 > 0$ and $\sigma_e^2 > 0$.
- The model also assumes that $\text{var}(Y_{h\ell j})$ is constant for all j , so that overall variance does not **change** over time.
- Finally, (3.6) shows that the model implies that $\text{var}(\mathbf{Y}_{h\ell})$ is assumed to be **the same** for all groups $\ell = 1, \dots, g$, which reflects the assumed independence of $b_{h\ell}$ and $\mathbf{e}_{h\ell}$ from the among-individual covariate group.

RESULT: This modeling approach and its assumptions induce a **compound symmetric** model for the overall aggregate pattern of correlation that is the same in each group. As we have noted previously, the compound symmetric model can be a restrictive representation of the overall pattern of correlation in the case **within-individual** sources of correlation are nonnegligible. The compound symmetric structure emphasizes **among-individual** sources of correlation, so may be reasonable when these sources are dominant.

The approach also induces the restriction that **overall variance** is **constant** across observation times. This is **may not** always be a realistic assumption for longitudinal data, as in many settings overall variance exhibits an **increasing** pattern over time.

Historically, the analysis methods we discuss shortly that are associated with this model have been used widely in agricultural, social science, and a host of other application areas, particularly before the advent of more modern methods, with little attention paid to the validity of these and other embedded assumptions. It is important that the data analyst understand the restrictions this approach involves.

ALTERNATIVE MODEL REPRESENTATION: It is of course possible to express model (3.4) in terms of the notation developed in Chapter 2. Recognizing that each (h, ℓ) , $h = 1, \dots, r_\ell$, $\ell = 1, \dots, g$, indexes one of $m = \sum_{\ell=1}^g r_\ell$ unique individuals, we can reindex individuals, and thus $\mathbf{Y}_{h\ell}$, $b_{h\ell}$, and $\mathbf{e}_{h\ell}$, and $\epsilon_{h\ell}$ in (3.5), using a single index $i = 1, \dots, m$ and reexpress the model in the form

$$\mathbf{Y}_i = \boldsymbol{\mu}_i + \mathbf{B}_i + \mathbf{e}_i$$

as in (2.9), where $\mathbf{Y}_i = (Y_{i1}, \dots, Y_{in})^T$ for individual i , as follows.

For illustration, take $g = 2$ and $n = 3$, and suppose individual i is in group 2. Then $E(\mathbf{Y}_i) = \boldsymbol{\mu}_i$ is such that $\boldsymbol{\mu}_i$ is equal to $\boldsymbol{\mu}_\ell$ in (3.5) with $\ell = 2$; that is,

$$\boldsymbol{\mu}_i = \begin{pmatrix} \mu_{21} \\ \mu_{22} \\ \mu_{23} \end{pmatrix} = \begin{pmatrix} \mu + \tau_2 + \gamma_1 + (\tau\gamma)_{21} \\ \mu + \tau_2 + \gamma_2 + (\tau\gamma)_{22} \\ \mu + \tau_2 + \gamma_3 + (\tau\gamma)_{23} \end{pmatrix},$$

and the model can be written as

$$\mathbf{Y}_i = \boldsymbol{\mu}_i + \mathbf{1}b_i + \mathbf{e}_i = \boldsymbol{\mu}_i + \epsilon_i. \quad (3.9)$$

In fact, defining

$$\boldsymbol{\beta} = (\mu, \tau_1, \tau_2, \gamma_1, \gamma_2, \gamma_3, (\tau\gamma)_{11}, (\tau\gamma)_{12}, (\tau\gamma)_{13}, (\tau\gamma)_{21}, (\tau\gamma)_{22}, (\tau\gamma)_{23})^T,$$

and

$$\mathbf{X}_i = \left(\begin{array}{c|c|c|c|c} 1 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \end{array} \right),$$

(3.9) can be written as

$$\mathbf{Y}_i = \mathbf{X}_i\boldsymbol{\beta} + \mathbf{1}b_i + \mathbf{e}_i \quad \text{or} \quad \mathbf{Y}_i = \mathbf{X}_i\boldsymbol{\beta} + \epsilon_i, \quad i = 1, \dots, m. \quad (3.10)$$

All of this of course generalizes to any g and n .

- In this notation, information on group membership for individual i is incorporated in the definition of μ_i and in particular, in (3.10), in the “design matrix” \mathbf{X}_i . All individuals with the same level of the group factor (i.e., sharing the same value of the among-individual covariate defining groups) have the same μ_i and \mathbf{X}_i .
- As in an analysis of variance formulation, the “design matrix” \mathbf{X}_i is **not of full rank**, so that the “model” for the overall population mean $\mu_i = \mathbf{X}_i\beta$ for any individual is **overparameterized**. To achieve a unique representation and **identify** the parameters μ , τ_ℓ , γ_j , and $(\tau\gamma)_{\ell j}$ for $\ell = 1, \dots, g$ and $j = 1, \dots, n$, it is customary to impose the following **constraints**:

$$\sum_{\ell=1}^g \tau_\ell = 0, \quad \sum_{j=1}^n \gamma_j = 0, \quad \sum_{\ell=1}^g (\tau\gamma)_{\ell j} = 0 = \sum_{j=1}^n (\tau\gamma)_{\ell j} \text{ for all } j, \ell, \quad (3.11)$$

which is equivalent to redefining the vector of parameters β and the matrices \mathbf{X}_i so that \mathbf{X}_i is of **full rank** for all i .

QUESTIONS OF INTEREST AND STATISTICAL HYPOTHESES: As we have noted, a common objective in the analysis of longitudinal data is to assess if the way in which response changes over time is different for different populations of individuals that can be distinguished by values of among-individual covariates like gender in the dental study or dose in the guinea pig study. In **classical** statistical analysis, such questions are interpreted as pertaining to **population mean response**; e.g., in the dental study, is **pattern of change of mean response** over age different for the populations of boys and girls?

Figure 3.8 depicts for $g = 2$ groups and $n = 3$ time points two situations in which the mean responses for each group for the three times lie on a **straight line**. In the left panel, the **rate of change**, represented by the **slope** of the two lines, is the **same** for both groups, so that the lines are **parallel**, whereas in the right panel the rate of change for group 2 is steeper than for group 1, so that the lines are not parallel. Thus, in the left panel, the pattern of change is **the same** while in the right it is **different**.

The left panel of Figure 3.9 shows for $g = 2$ groups and $n = 3$ time points a scenario where the **mean response profiles** are also **parallel**, so that the pattern of change in each group is the same, but the means at each time **do not** follow an apparent straight line relationship. The right panel shows a case where the pattern is different.

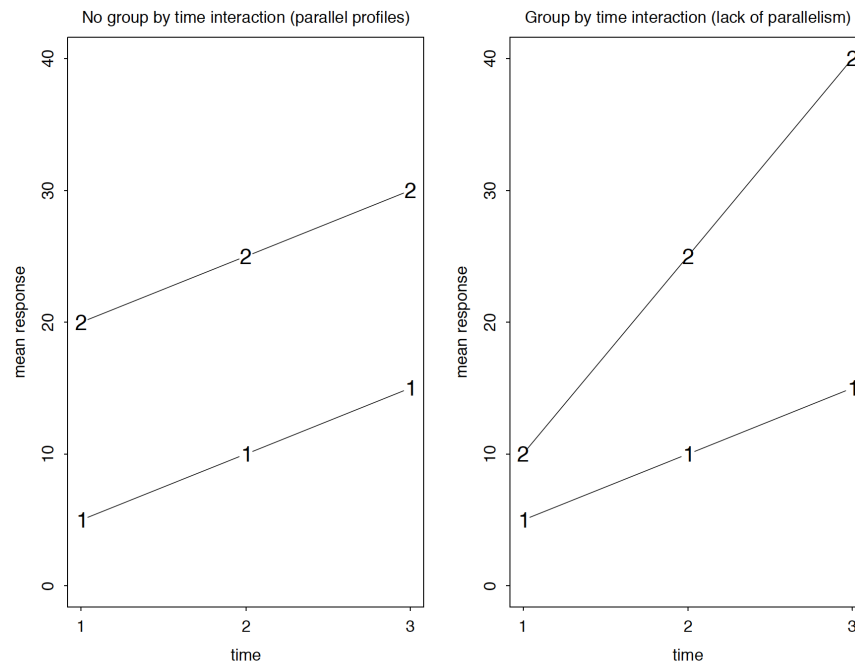


Figure 3.8: *Straight line mean profiles. Mean response for each group at each time, where the plotting symbol indicates group number. There is no interaction in the left panel; the right panel shows a quantitative interaction.*

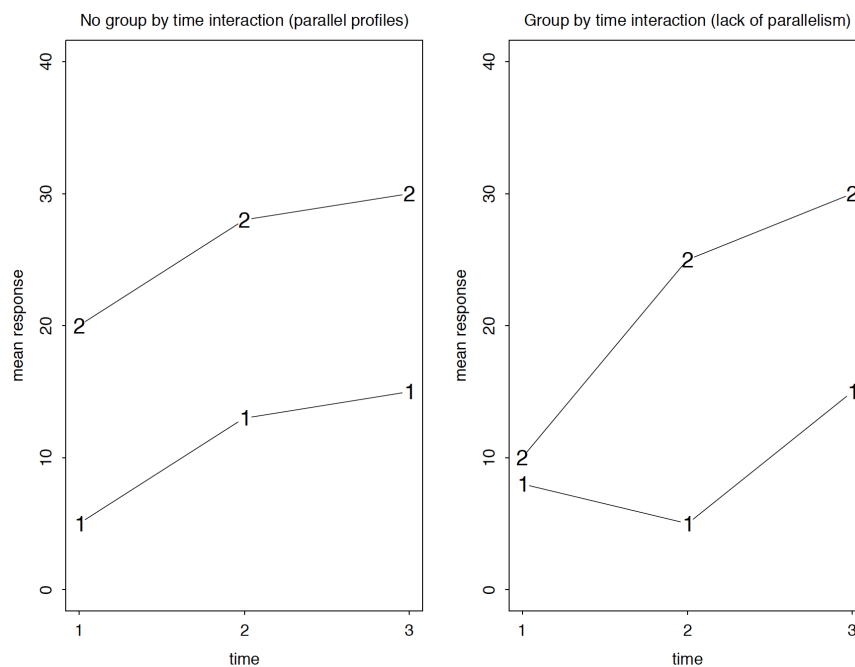


Figure 3.9: *Mean profiles not a straight line.*

GROUP BY TIME INTERACTION: In classical jargon, the situations in the right hand panels of Figures 3.8 and 3.9 depict examples of a **group by time interaction**; in each panel, the difference in mean response between groups is **not the same** at all time points. In both figures, the **direction** of the difference in mean response between groups is **the same** in that the mean for group 2 is always larger. This is often referred to as a **quantitative interaction**, particularly in health sciences research.




Figure 3.10: *Crossed mean profiles. Each panel represents a form of qualitative interaction.*

Figure 3.10 shows a different type of group by time interaction. In both panels, the difference in mean response between groups is again not the same at all time points, and the **direction** of the difference is **not the same**, either. For example, in the left hand panel, the **magnitude** of the difference at times 1 and 3 is the same, but in the opposite direction. This is often referred to as a **qualitative interaction**.

Returning to the model expressed using the classical notation as in (3.1), each mean in all of these figures is represented by

$$\mu_{\ell j} = \mu + \tau_{\ell} + \gamma_j + (\tau\gamma)_{\ell j}.$$

The difference between mean response in groups 1 and 2 at any time j is, under this model,

$$\mu_{1j} - \mu_{2j} = (\tau_1 - \tau_2) + \{(\tau\gamma)_{1j} - (\tau\gamma)_{2j}\}.$$

Thus, the $(\tau\gamma)_{\ell j}$ allow the difference in means between groups to be **different** at different times j , as in the right panels of Figures 3.8 and 3.9 and in Figure 3.10, by the amount $\{(\tau\gamma)_{1j} - (\tau\gamma)_{2j}\}$ at time j .

If the $(\tau\gamma)_{\ell j}$ were all the **same**, the difference in means at any j reduces to

$$\mu_{1j} - \mu_{2j} = (\tau_1 - \tau_2),$$

so that the difference in mean response between groups is **the same at all time points** and equal to $(\tau_1 - \tau_2)$, which does not depend on j . This is the case in the left panels of Figures 3.8 and 3.9. Here, the **pattern of change** over time is thus **the same**; i.e., the mean profiles for each group are **parallel** over time.

Under the constraints

$$\sum_{\ell=1}^g (\tau\gamma)_{\ell j} = 0 = \sum_{j=1}^n (\tau\gamma)_{\ell j} \text{ for all } \ell, j$$

in (3.11), if $(\tau\gamma)_{\ell j}$ are all **the same** for all ℓ, j , then it must be that

$$(\tau\gamma)_{\ell j} = 0 \text{ for all } \ell, j.$$

In general, then, if we wish to address the question of whether or not there is a common pattern of change over time, so whether or not the mean profiles are **parallel**, we can cast this in terms of the **null hypothesis**

$$H_0 : \text{all } (\tau\gamma)_{\ell j} = 0, \quad \ell = 1, \dots, g, \quad j = 1, \dots, n, \quad (3.12)$$

with the alternative being that at least one $(\tau\gamma)_{\ell j} \neq 0$, in which case the mean difference at at least one of the time points is different from that at the others.

There are gn parameters $(\tau\gamma)_{\ell j}$; however, if the constraints above hold, then having $(g-1)(n-1)$ of the $(\tau\gamma)_{\ell j}$ equal to 0 automatically requires the remaining ones to be zero. Thus, the hypothesis (3.12) is really one about the behavior of $(g-1)(n-1)$ parameters, so there are $(g-1)(n-1)$ **degrees of freedom** associated with this hypothesis.

In the classical literature on analysis of variance for repeated measurements, the test of the null hypothesis (3.12) is referred to as the **test for parallelism**. As Figure 3.9 demonstrates, parallelism does not necessarily mean that the pattern of mean response in each group follow a **straight line**.

MAIN EFFECT OF GROUPS: If mean profiles are parallel, then the obvious next question is whether or not they are **coincident**; that is, whether or not the mean response is in fact **the same** for each group at each time point. A little thought reveals that, if the mean profiles are **parallel**, if they are furthermore coincident, then the **average** of the mean responses over time will be the same for each group. The question of whether or not the average of mean responses is the same for each group if the profiles are **not parallel** may or may not be interesting or relevant.

- If in truth the situation were like those depicted in the right hand panels of Figures 3.8 and 3.9, whether or not the average of mean responses over time is different for the two groups might be interesting, as it would reflect that the mean response for group 2 is larger **at all times**.
- On the other hand, consider the left panel of Figure 3.10. If this were the true state of affairs, this issue is **meaningless**; the change of mean response over time is in the **opposite** direction for the two groups; thus, how it averages out over time is of little importance. Because the phenomenon of interest does indeed happen **over time**, the **average** of what it does over time may be something that cannot be achieved – we can't make time stand still.
- Similarly, if the issue under study is something like growth, the **average** over time of the response may have little meaning; instead, one may be interested in, for example, how different the mean response is at the end of the time period of study. For example, in the right panel of Figure 3.10, mean response over time increases for each group at different rates, but has the same average over time. The group with the steeper rate will have a larger mean response at the end of the time period.

In general, then, the question of whether or not the average of the mean response over time is the same across groups in a longitudinal study is of most interest when the mean profiles over time are approximately parallel.

For definiteness, consider the case of $g = 2$ groups and $n = 3$ time points. For group ℓ , the average of means over time is, with $n = 3$,

$$n^{-1}(\mu_{\ell 1} + \mu_{\ell 2} + \mu_{\ell 3}) = \mu + \tau_{\ell} + n^{-1}(\gamma_1 + \gamma_2 + \gamma_3) + n^{-1}\{(\tau\gamma)_{\ell 1} + (\tau\gamma)_{\ell 2} + (\tau\gamma)_{\ell 3}\}.$$

The difference of the averages between $\ell = 1$ and $\ell = 2$ is then (algebra)

$$\tau_1 - \tau_2 + n^{-1} \sum_{j=1}^n (\tau\gamma)_{1j} - n^{-1} \sum_{j=1}^n (\tau\gamma)_{2j}.$$

The **constraints** (3.11) imposed to render the model of **full rank** dictate that $\sum_{j=1}^n (\tau\gamma)_{\ell j} = 0$ for each ℓ ; thus, the two sums in this expression are 0 by assumption, so that we are left with $\tau_1 - \tau_2$.

Thus, the hypothesis may be expressed as

$$H_0 : \tau_1 - \tau_2 = 0.$$

Furthermore, under the constraint $\sum_{\ell=1}^g \tau_{\ell} = 0$, if the τ_{ℓ} are equal as in H_0 , then they must satisfy $\tau_{\ell} = 0$ for each ℓ . Thus, the hypothesis may be rewritten as

$$H_0 : \tau_1 = \tau_2 = 0.$$

For general g and n , the reasoning is the same; we have

$$H_0 : \tau_1 = \dots = \tau_g = 0. \quad (3.13)$$

MAIN EFFECT OF TIME: A further question of interest may be whether or not the mean response is in fact **constant** over time. If the profiles are parallel, then this is like asking whether the mean response averaged across groups is the **same** at each time. If the profiles are not parallel, then this may or may not be interesting. For example, in the left panel of Figure 3.10, the average of mean responses for groups 1 and 2 are the same at each time point. However, the mean response is certainly not constant across time for either group. If the groups represent a factor like gender, then what happens on average is something that can never be achieved.

The average of mean responses across groups for time j is

$$g^{-1} \sum_{\ell=1}^g \mu_{\ell j} = \gamma_j + q^{-1} \sum_{\ell=1}^g \tau_{\ell} + q^{-1} \sum_{\ell=1}^g (\tau\gamma)_{\ell j} = \gamma_j$$

using the constraints $\sum_{\ell=1}^g \tau_{\ell} = 0$ and $\sum_{\ell=1}^g (\tau\gamma)_{\ell j} = 0$ in (3.11). Thus, in the special case $g = 2$ and $n = 3$, having all these averages be the same at each time is equivalent to

$$H_0 : \gamma_1 = \gamma_2 = \gamma_3.$$

Under the constraint $\sum_{j=1}^n \gamma_j = 0$, then, we have $H_0 : \gamma_1 = \gamma_2 = \gamma_3 = 0$. For general g and n , the hypothesis is of the form

$$H_0 : \gamma_1 = \dots = \gamma_n = 0. \quad (3.14)$$

REMARK: Hypotheses (3.12), (3.13), and (3.14) are, of course, exactly the hypotheses that one tests for a **split plot experiment**, where, here, “time” plays the role of the “split plot” factor and “group” is the “whole plot factor.” What is different is the **interpretation**; because “time” has a natural **ordering** (longitudinal), what is interesting may be different; as noted above, of primary interest is whether or not the pattern of change in mean response over levels of time is different across groups.

ANALYSIS OF VARIANCE: Given that the statistical model and hypotheses of interest here are **identical** to those for a split plot, it should come as no surprise that the analysis is identical. Under the assumption that the model (3.1) is correctly specified and that the responses are normally distributed, so that

$$\mathbf{Y}_{h\ell} \sim \mathcal{N}_n(\boldsymbol{\mu}_\ell, \mathbf{V}), \quad \mathbf{V} = \sigma_b^2 \mathbf{J}_n + \sigma_e^2 \mathbf{I}_n. \quad (3.15)$$

as in (3.6), it can be shown the **F ratios** one would construct under the usual principles of analysis of variance provide the basis for valid tests of the hypotheses above. For brevity, we present the **analysis of variance table** and associated testing procedures without proof.

Define

- $\bar{Y}_{h\ell} = n^{-1} \sum_{j=1}^n Y_{h\ell j}$, the sample average over time for the h th unit in the ℓ th group (over all observations on this unit)
- $\bar{Y}_{\cdot\ell j} = r_\ell^{-1} \sum_{h=1}^{r_\ell} Y_{h\ell j}$, the sample average at time j in group ℓ over all units
- $\bar{Y}_{\cdot\ell} = (r_\ell n)^{-1} \sum_{h=1}^{r_\ell} \sum_{j=1}^n Y_{h\ell j}$, the sample average of all observations in group ℓ
- $\bar{Y}_{\cdot j} = m^{-1} \sum_{\ell=1}^g \sum_{h=1}^{r_\ell} Y_{h\ell j}$, the sample average of all observations at the j th time
- $\bar{Y}_{\dots} =$ the average of all mn observations.

Let

$$\begin{aligned} SS_G &= \sum_{\ell=1}^g nr_\ell (\bar{Y}_{\cdot\ell} - \bar{Y}_{\dots})^2, \quad SS_{Tot,U} = n \sum_{\ell=1}^g \sum_{h=1}^{r_\ell} (\bar{Y}_{h\ell} - \bar{Y}_{\dots})^2 \\ SS_T &= m \sum_{j=1}^n (\bar{Y}_{\cdot j} - \bar{Y}_{\dots})^2, \quad SS_{GT} = \sum_{j=1}^n \sum_{\ell=1}^g r_\ell (\bar{Y}_{\cdot\ell j} - \bar{Y}_{\dots})^2 - SS_T - SS_G \\ SS_{Tot,all} &= \sum_{\ell=1}^g \sum_{h=1}^{r_\ell} \sum_{j=1}^n (Y_{h\ell j} - \bar{Y}_{\dots})^2. \end{aligned}$$

Then the following analysis of variance table is constructed.

Source	SS	DF	MS	F
Among Groups	SS_G	$g - 1$	MS_G	$F_G = MS_G / MS_{EU}$
Among-Unit Error	$SS_{Tot,U} - SS_G$	$m - g$	MS_{EU}	
Time	SS_T	$n - 1$	MS_T	$F_T = MS_T / MS_E$
Group \times Time	SS_{GT}	$(g - 1)(n - 1)$	MS_{GT}	$F_{GT} = MS_{GT} / MS_E$
Within-Unit Error	SS_E	$(m - g)(n - 1)$	MS_E	
Total	$SS_{Tot,all}$	$nm - 1$		

Here, the **mean squares** (MS) for each source are equal to the sum of squares (SS) divided by the degrees of freedom; e.g., $MS_G = SS_G / (g - 1)$, and

$$SS_E = SS_{Tot,all} - SS_{GT} - SS_T - SS_{Tot,U}.$$

REMARK: It is traditional in the classical terminology to use the term “error;” however, it is important to recognize that the “Among Unit Error” includes variation due **among-individual variability** and the “Within-Unit Error” includes variation due to both **within-individual realization (fluctuations)** and **measurement error**.

Under (3.15), the **expectations** of the mean squares in the table can be derived; we do not present these calculations here (they can be found in Section 3.3 of Crowder and Hand, 1990). These **expected mean squares** are shown in the following table; these are valid only if the model (3.15) is indeed correctly specified, so that the true overall, aggregate pattern of correlation is **compound symmetric**.

Source	MS	Expected mean square
Among Groups	MS_G	$\sigma_e^2 + n\sigma_b^2 + n \sum_{\ell=1}^g r_\ell \tau_\ell^2 / (g - 1)$
Among-Unit error	MS_{EU}	$\sigma_e^2 + n\sigma_b^2$
Time	MS_T	$\sigma_e^2 + m \sum_{j=1}^n \gamma_j^2 / (n - 1)$
Group \times Time	MS_{GT}	$\sigma_e^2 + \sum_{\ell=1}^g r_\ell \sum_{j=1}^n (\tau\gamma)_{\ell j}^2 / (g - 1)(n - 1)$
Within-Unit Error	MS_E	σ_e^2

Inspection of the expected mean squares shows informally that we expect the F ratios in the analysis of variance table to test the appropriate issues. For example, we would expect F_{GT} to be large if the $(\tau\gamma)_{\ell j}$ are not all equal to zero, and F_G to be large if the τ_ℓ are not all equal to zero.

Note that F_G uses the appropriate denominator; intuitively, we wish to compare the mean square for groups against an “error term” that takes into account **all** sources of variation **among** (σ_b^2) and **within** (σ_e^2) individuals. The other two tests are on features that occur **within individuals**; thus, the denominator takes account of the relevant source of variation, that within individuals (σ_e^2).

It can be shown formally that, as long as (3.15) is correctly specified, under the null hypotheses (3.12), (3.13), and (3.14), the **sampling distributions** of the F ratios in the analysis of variance table are F distributions with the degrees of freedom specified below.

TEST PROCEDURES: We now summarize the procedures for testing each of the hypotheses. Here, $\mathcal{F}_{a,b,\alpha}$ is the critical value corresponding to level of significance α for an F distribution with a numerator and b denominator degrees of freedom.

- **Group by time interaction (parallelism), (3.12).**

$$H_0 : (\tau\gamma)_{\ell j} = 0 \text{ for all } j, \ell \text{ vs. } H_1 : \text{ at least one } (\tau\gamma)_{\ell j} \neq 0.$$

A valid test rejects H_0 at level of significance α if

$$F_{GT} > \mathcal{F}_{(g-1)(n-1), (n-1)(m-g), \alpha}$$

or, equivalently, if the probability is less than α that one would see a value of the test statistic as large or larger than F_{GT} if H_0 were true (that is, the p-value is less than α).

- **Main effect of group (coincidence), (3.13).**

$$H_0 : \tau_\ell = 0 \text{ for all } \ell \text{ vs. } H_1 : \text{ at least one } \tau_\ell \neq 0.$$

A valid test rejects H_0 at level of significance α if

$$F_G > \mathcal{F}_{g-1, m-g, \alpha}.$$

- **Main effect of time (constancy), (3.14).**

$$H_0 : \gamma_j = 0 \text{ for all } j \text{ vs. } H_1 : \text{ at least one } \gamma_j \neq 0.$$

A valid test rejects H_0 at level α if

$$F_T > \mathcal{F}_{n-1, (n-1)(m-g), \alpha}.$$

VIOLATION OF COVARIANCE MATRIX ASSUMPTION: These test procedures are valid under model (3.15), which embodies the assumption of **compound symmetry** of the overall correlation matrix of a data vector. In fact, it can be shown that they are also valid under slightly **more general conditions** that include compound symmetry as a special case. However, validity of the tests is predicated on the covariance matrix being of the special form we discuss next; if not, then F ratios F_T and F_{GT} **no longer** have exactly an F distribution, and the associated tests are not valid and can lead to erroneous conclusions.

A $(n \times n)$ matrix \mathbf{V} is said to be of **Type H** if it can be written in the form

$$\mathbf{V} = \begin{pmatrix} \lambda + 2\alpha_1 & \alpha_1 + \alpha_2 & \cdots & \alpha_1 + \alpha_n \\ \alpha_2 + \alpha_1 & \lambda + 2\alpha_2 & \cdots & \alpha_2 + \alpha_n \\ \vdots & \vdots & \ddots & \vdots \\ \alpha_n + \alpha_1 & \alpha_n + \alpha_2 & \cdots & \lambda + 2\alpha_n \end{pmatrix}. \quad (3.16)$$

It is straightforward to deduce that a covariance matrix with correlation structure that is compound symmetric is of Type H.

It can be shown that, as long as the data vectors $\mathbf{Y}_{h\ell}$ are multivariate normal with common covariance matrix \mathbf{V} of the form (3.16), the F tests discussed above **will be valid**. Thus, because (3.16) includes compound symmetry, the tests are valid if model 3.15) holds. If the overall covariance matrix is **not** of Type H, but these F tests are conducted nonetheless, they will be too **liberal**; that is, they will reject the null hypothesis more often than they should, so that, for example, the analyst might conclude that there is sufficient evidence supporting a group by time interaction when there really is not.

It is possible to construct tests of whether or not the true overall covariance matrix is of Type H. One such test is **Mauchly's test for sphericity**. We do not present the form and derivation of this test here; description of the test is given by Vonesh and Chinchilli (1997, p. 85), for example. The test statistic for testing the null hypothesis

$$H_0 : \mathbf{V} \text{ is of Type H,}$$

where \mathbf{V} is the true covariance matrix of a response vector, has approximately a χ^2 (chi-square) distribution when the number of individuals m is "large," with degrees of freedom $(n - 2)(n + 1)/2$. Thus, the test is performed at level of significance α by comparing the value of the test statistic to the χ^2_{α} critical value with $(n - 2)(n + 1)/2$ degrees of freedom.

All such tests have **limitations**: they are not very powerful with the numbers of individuals r_ℓ in each group is not large, and they can be misleading if the true distribution of the response vectors is not multivariate normal. Accordingly, we do not discuss it further. We return to the issue of approaches when the analyst lacks confidence in the validity of the assumption of Type H covariance structure in the next section.

3.3 Specialized within-individual hypotheses and tests

The hypotheses of group by time interaction (parallelism) and main effect of time (constancy) have to do with questions about the **pattern of change** over time. However, they address these issues in an “overall” sense; e.g., the test of the group by time interaction asks only if the pattern of mean responses over time is different for different groups, but it does **not provide insight** into the nature of the pattern of change and how it differs.

We now review methods to carry out a **more detailed study** of specific aspects of how the mean response changes over time. As we demonstrate, these methods do this through testing of specialized hypotheses.

It is conventional to present the relevant null hypotheses, and, indeed, the three main null hypotheses (3.12), (3.13), and (3.14) using the following **unified notation**. Let \mathcal{M} denote the matrix of all means $\mu_{\ell j}$ implied by the model (3.1), i.e.

$$\mathcal{M} = \begin{pmatrix} \mu_{11} & \mu_{12} & \cdots & \mu_{1n} \\ \vdots & \vdots & \vdots & \vdots \\ \mu_{g1} & \mu_{g2} & \cdots & \mu_{gn} \end{pmatrix} = \begin{pmatrix} \mu_1^T \\ \vdots \\ \mu_g^T \end{pmatrix}, \quad (3.17)$$

so that the ℓ th row of \mathcal{M} in (3.17) is μ_ℓ^T . Let

- \mathbf{C} be a $(c \times g)$ matrix with $c \leq g$ of full rank.
- \mathbf{U} be a $(n \times u)$ matrix with $u \leq n$ of full rank.

Then it is possible to express null hypotheses of interest in the **general form**

$$H_0 : \mathbf{CMU} = \mathbf{0}. \quad (3.18)$$

In this formulation

- the matrix \mathbf{C} specifies **differences among** or **averages across** groups
- the matrix \mathbf{U} specifies **differences over** or **averages across** levels of time.

Depending on the choices of the matrices \mathbf{C} and \mathbf{U} in (3.18), the resulting **linear function** \mathbf{CMU} of the elements of \mathcal{M} (the individual means for different groups at different time points) can be made to address **specialized questions** regarding differences in mean response among groups and in patterns of change over time.

To see this, first consider the null hypothesis for the **group by time interaction** (parallelism) (3.12), $H_0 : \text{all } (\tau\gamma)_{\ell j} = 0$, with $g = 2$ groups and $n = 3$ time points. Take

$$\mathbf{C} = \begin{pmatrix} 1, & -1 \end{pmatrix}, \quad (3.19)$$

so that $c = 1 = q - 1$. Note that

$$\begin{aligned} \mathbf{CM} &= \begin{pmatrix} 1, & -1 \end{pmatrix} \begin{pmatrix} \mu_{11} & \mu_{12} & \mu_{13} \\ \mu_{21} & \mu_{22} & \mu_{23} \end{pmatrix} = \begin{pmatrix} \mu_{11} - \mu_{21}, & \mu_{12} - \mu_{22}, & \mu_{13} - \mu_{23} \end{pmatrix} \\ &= \begin{pmatrix} \tau_1 - \tau_2 + (\tau\gamma)_{11} - (\tau\gamma)_{21}, & \tau_1 - \tau_2 + (\tau\gamma)_{12} - (\tau\gamma)_{22}, & \tau_1 - \tau_2 + (\tau\gamma)_{13} - (\tau\gamma)_{23} \end{pmatrix} \end{aligned}$$

Thus, \mathbf{C} yields differences in means among groups at each time point.

Take

$$\mathbf{U} = \begin{pmatrix} 1 & 0 \\ -1 & 1 \\ 0 & -1 \end{pmatrix}, \quad (3.20)$$

so that $u = 2 = n - 1$. Thus, \mathbf{U} involves differences for pairs of time points.

It is straightforward (try it) to show that

$$\begin{aligned} \mathbf{CMU} &= \begin{pmatrix} \mu_{11} - \mu_{21} - \mu_{12} + \mu_{22}, & \mu_{12} - \mu_{22} - \mu_{13} + \mu_{23} \end{pmatrix} \\ &= \begin{pmatrix} (\tau\gamma)_{11} - (\tau\gamma)_{21} - (\tau\gamma)_{12} + (\tau\gamma)_{22}, & (\tau\gamma)_{12} - (\tau\gamma)_{22} - (\tau\gamma)_{13} + (\tau\gamma)_{23} \end{pmatrix}. \end{aligned}$$

It is an exercise in algebra to verify that, under the constraints in (3.11), if each of these elements equals zero, then H_0 follows.

Similarly, for the null hypothesis for the **main effect of groups** (coincidence), $H_0 : \tau_1 = \tau_2 = 0$, taking \mathbf{C} to be as in (3.19) and

$$\mathbf{U} = \begin{pmatrix} 1/3 \\ 1/3 \\ 1/3 \end{pmatrix},$$

it is straightforward to see that, with $n = 3$,

$$\mathbf{CMU} = \tau_1 - \tau_2 + n^{-1} \sum_{j=1}^n (\tau\gamma)_{1j} - n^{-1} \sum_{j=1}^n (\tau\gamma)_{2j}.$$

That is, this choice of \mathbf{U} dictates an **averaging** operation across time. Imposing the constraints as above, we can express H_0 in the form $H_0 : \mathbf{CMU} = 0$.

To express the null hypothesis for the **main effect of time** (constancy), $H_0 : \gamma_1 = \gamma_2 = \gamma_3 = 0$, take

$$\mathbf{U} = \begin{pmatrix} 1 & 0 \\ -1 & 1 \\ 0 & -1 \end{pmatrix}, \quad \mathbf{C} = \begin{pmatrix} 1/2, & 1/2 \end{pmatrix}.$$

Here, \mathbf{C} involves **averaging across groups** while \mathbf{U} involves differences for pairs of time points as above. Then

$$\begin{aligned} \mathbf{MU} &= \begin{pmatrix} \mu_{11} & \mu_{12} & \mu_{13} \\ \mu_{21} & \mu_{22} & \mu_{23} \end{pmatrix} \begin{pmatrix} 1 & 0 \\ -1 & 1 \\ 0 & -1 \end{pmatrix} = \begin{pmatrix} \mu_{11} - \mu_{12} & \mu_{12} - \mu_{13} \\ \mu_{21} - \mu_{22} & \mu_{22} - \mu_{23} \end{pmatrix} \\ &= \begin{pmatrix} \gamma_1 - \gamma_2 + (\tau\gamma)_{11} - (\tau\gamma)_{12}, & \gamma_2 - \gamma_3 + (\tau\gamma)_{12} - (\tau\gamma)_{13} \\ \gamma_1 - \gamma_2 + (\tau\gamma)_{21} - (\tau\gamma)_{22}, & \gamma_2 - \gamma_3 + (\tau\gamma)_{22} - (\tau\gamma)_{23} \end{pmatrix}. \end{aligned} \quad (3.21)$$

from whence it is straightforward to derive, imposing the constraints in (3.11), that

$$\mathbf{CMU} = \begin{pmatrix} \gamma_1 - \gamma_2, & \gamma_2 - \gamma_3 \end{pmatrix}.$$

Setting this equal to zero with the constraint $\sum_{j=1}^n \gamma_j = 0$ yields H_0 .

Clearly, the principles involved in specifying the matrices \mathbf{C} and \mathbf{U} to yield the form of the null hypotheses corresponding to the group by time interaction, main effect of groups, and main effect of time generalize to any g and n .

Other choices of \mathbf{C} and \mathbf{U} can be made to examine **components** making up these overall hypotheses and to isolate **specific features** of the pattern of change. Recall the following definition.

CONTRAST: If \mathbf{c} is a $(n \times 1)$ vector and $\boldsymbol{\mu}$ is a $(n \times 1)$ vector of means, then the **linear combination** $\mathbf{c}^T \boldsymbol{\mu} = \boldsymbol{\mu}^T \mathbf{c}$ is a **contrast** if \mathbf{c} is such that its elements sum to zero; i.e., $\mathbf{c}^T \mathbf{1} = 0$.

Thus, for example, with $g = 2$ and $n = 3$, the columns of the matrix

$$\mathbf{U} = \begin{pmatrix} 1 & 0 \\ -1 & 1 \\ 0 & -1 \end{pmatrix}$$

in (3.20) define contrasts of elements of $\boldsymbol{\mu}_1$ and $\boldsymbol{\mu}_2$, the mean vectors for groups 1 and 2, in that

$$\mathcal{M}\mathbf{U} = \begin{pmatrix} \mu_{11} - \mu_{12} & \mu_{12} - \mu_{13} \\ \mu_{21} - \mu_{22} & \mu_{22} - \mu_{23} \end{pmatrix}. \quad (3.22)$$

In (3.22), each entry is a **contrast** involving differences in mean response between pairs of times in each group. Specialized questions of interest can be posed by considering these contrasts.

- The **difference** of the contrasts in the first column of (3.22) focuses on whether or not the way in which the mean response differs from time 1 to time 2 is different in groups 1 and 2. This feature is clearly a **component** of the overall group by time interaction, focusing in particular on times 1 and 2. Likewise, the difference of the contrasts in the second columns of (3.22) focuses on the same for times 2 and 3, and is also part of the group by time interaction.

Indeed, taken together, the differences of contrasts in both columns of (3.22) **fully characterize** the overall group by time interaction.

- Similarly, the **average** of the contrasts in the first column of (3.22) focuses on the difference in mean response between times 1 and 2, averaged across groups. This is clearly a component of the main effect of time. Similarly, the average of contrasts in the second column reflects the same for times 2 and 3. Again, taken together, the averages of contrasts in both rows of (3.22) **fully characterize** the overall main effect of time.

Thus, considering these contrasts and their differences among or averages across groups serves to “**pick apart**” how the overall group by time interaction effect and main effect of time occur and can provide insight into **specific features** of the pattern of change over time.

PROFILE TRANSFORMATION: For general number of groups g and number of time points n , the extension of (3.20) is the $(n \times n - 1)$ matrix

$$\mathbf{U} = \begin{pmatrix} 1 & 0 & \cdots & 0 \\ -1 & 1 & \cdots & 0 \\ 0 & -1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & \cdots & \cdots & 1 \\ 0 & \cdots & 0 & -1 \end{pmatrix}. \quad (3.23)$$

Postmultiplication of \mathcal{M} by \mathbf{U} in (3.23) results in contrasts comparing means at successive pairs of time points and is often called the **profile transformation** of the means. Examining individually the differences among or averages across the contrasts resulting from each column provides insight on the contribution to the overall pattern of change over time.

Other \mathbf{U} matrices allow other ways of “parsing” the pattern of change over time. For example, instead of focusing on changes from one time to the next, one might consider how the mean at a specific time point differs from what happens at **all subsequent time points**. This might highlight at what point in time changes in the pattern begin to emerge.

We demonstrate with $g = 2$ and $n = 4$ and consider the contrast

$$\mu_{11} - (\mu_{12} + \mu_{13} + \mu_{14})/3,$$

which compares, for group 1, the mean at time 1 to the **average** of means at all other times. Similarly,

$$\mu_{12} - (\mu_{13} + \mu_{14})/2$$

compares for group 1 the mean at time 2 to the average of those at subsequent times. The final contrast of this type for group 1 is

$$\mu_{13} - \mu_{14},$$

which compares what happens at time 3 to the “average” of what comes next, the single mean at time 4. We may similarly specify such contrasts for the other group.

These contrasts can be obtained by postmultiplying \mathcal{M} by

$$\mathbf{U} = \begin{pmatrix} 1 & 0 & 0 \\ -1/3 & 1 & 0 \\ -1/3 & -1/2 & 1 \\ -1/3 & -1/2 & -1 \end{pmatrix}. \quad (3.24)$$

In particular, with $g = 2$,

$$\mathcal{M}\mathbf{U} = \begin{pmatrix} \mu_{11} - \mu_{12}/3 - \mu_{13}/3 - \mu_{14}/3, & \mu_{12} - \mu_{13}/2 - \mu_{14}/2, & \mu_{13} - \mu_{14} \\ \mu_{21} - \mu_{22}/3 - \mu_{23}/3 - \mu_{24}/3, & \mu_{22} - \mu_{23}/2 - \mu_{24}/2, & \mu_{23} - \mu_{24} \end{pmatrix}. \quad (3.25)$$

HELMERT TRANSFORMATION: For general n , the $(n \times n-1)$ matrix whose columns define contrasts of this type is the so-called **Helmert transformation** matrix of the form

$$\mathbf{U} = \begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ -1/(n-1) & 1 & 0 & \cdots & 0 \\ -1/(n-1) & -1/(n-2) & 1 & \cdots & 0 \\ \vdots & \vdots & -1/(n-3) & \vdots & \vdots \\ -1/(n-1) & -1/(n-2) & \vdots & \cdots & 1 \\ -1/(n-1) & -1/(n-2) & -1/(n-3) & \cdots & -1 \end{pmatrix}. \quad (3.26)$$

Postmultiplication of \mathcal{M} by a matrix of the form (3.26) yields contrasts representing comparisons of each mean against the **average** of means at all subsequent times.

It is in fact the case that any \mathbf{U} matrix with $n-1$ columns, so involving $n-1$ contrasts that “pick apart” all possible differences in means over time, as do the **profile** and **Helmert** transformation matrices (3.23) and (3.26), lead to the overall hypotheses for group by time interaction and main effect of time when paired with the appropriate \mathbf{C} matrix.

For example, in the case $g = 2$, $n = 2$, if we premultiply **either** of (3.21) or (3.25) by $\mathbf{C} = (1/2, 1/2)$ and impose the constraints (3.11), we are led to the null hypothesis $H_0 : \gamma_1 = \gamma_2 = \gamma_3 = 0$ for the **main effect of time**.

In particular, with (3.21),

$$\mathbf{CM} \begin{pmatrix} 1 & 0 \\ -1 & 1 \\ 0 & -1 \end{pmatrix} = \begin{pmatrix} \gamma_1 - \gamma_2 & \gamma_2 - \gamma_3 \end{pmatrix} = \mathbf{0},$$

while with (3.25),

$$\mathbf{CM} \begin{pmatrix} 1 & 0 \\ -1/2 & 1 \\ -1/2 & -1 \end{pmatrix} = \begin{pmatrix} \gamma_1 - 0.5\gamma_2 - 0.5\gamma_3 & \gamma_2 - \gamma_3 \end{pmatrix} = \mathbf{0},$$

both of which can be shown to imply $\gamma_1 = \gamma_2 = \gamma_3$. The diligent student can verify a similar result for the group by time interaction with $\mathbf{C} = (1, -1)$.

In general, for a $(n \times n-1)$ \mathbf{U} matrix involving $n-1$ contrasts that characterizes all possible differences in means over time in a particular way,

- premultiplying \mathbf{MU} by the $(g-1) \times g$ matrix

$$\mathbf{C} = \begin{pmatrix} 1 & -1 & 0 & \cdots & 0 \\ 1 & 0 & -1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & 0 & 0 & \cdots & -1 \end{pmatrix},$$

considers how each contrast defined by the columns of \mathbf{U} differs across groups, which is a component of the **group by time interaction** (how the difference in means across groups is different at different times).

- premultiplying by $\mathbf{C} = (1/g, 1/g, \dots, 1/g)$, each of the $n-1$ elements of the resulting $1 \times (n-1)$ matrix correspond to the **average** of each of these contrasts over groups, which all together constitute the **main effect of time**.

SEPARATE COMPONENT TESTS: We can examine each of these components **separately** to explore particular aspects of the pattern of mean response over time. Formally, we can carry out separate hypothesis tests corresponding to each component. This can be accomplished as follows.

Consider the k th column of \mathbf{U} , denoted as \mathbf{c}_k , $k = 1, \dots, n - 1$.

- For each individual indexed by (h, ℓ) , obtain

$$\mathbf{c}_k' \mathbf{Y}_{h\ell},$$

which distills the repeated measurements on the individual to a single quantity representing the value of the component contrast for that individual. If $\text{var}(\mathbf{Y}_{h\ell}) = \mathbf{V}$, so is the same for all individuals, then these values have the **same variance** for all individuals (see below).

- Carry out analyses on the resulting “data;” e.g., to test if the component differs across groups, conduct a usual one-way analysis of variance on these “data” and carry out the F -test for the Group effect.
- To test if the contrast is zero averaged across groups, test whether the overall mean of the “data” is equal to zero using a standard t test (or equivalently, the F test based on the square of the t statistic).
- These tests will be valid **regardless** of whether or not **compound symmetry** holds; all that matters is that \mathbf{V} , whatever it is, is **the same** for all units, in which case

$$\text{var}(\mathbf{c}_k' \mathbf{Y}_{h\ell}) = \mathbf{c}_k' \mathbf{V} \mathbf{c}_k,$$

which is a constant for all h and ℓ , so that the usual assumption of constant variance necessary for the above analyses holds for the “data” corresponding to each contrast.

ORTHOGONAL CONTRASTS: Recall that if \mathbf{c}_1 and \mathbf{c}_2 are any two columns of \mathbf{U} , then if $\mathbf{c}_1' \mathbf{c}_2 = 0$, \mathbf{c}_1 and \mathbf{c}_2 are said to be **orthogonal**. The **contrasts** corresponding to these vectors are **orthogonal contrasts**. The contrasts corresponding to the columns of the **profile transformation** matrix (3.23) **are not** orthogonal, while those of the **Helmert transformation** matrix (3.26) **are** (try it).

There is an advantage to a transformation whose columns and thus embedded contrasts are orthogonal. As intuition might suggest, it can be shown that a set of $n - 1$ orthogonal contrasts **partitions** the total Group \times Time and Within-Unit Error sums of squares into $n - 1$ **distinct** or “nonoverlapping” components. Informally, this implies that the outcome of one of the component hypothesis tests addressing a particular contrast can be considered separately, regardless of the outcome of the tests for the others. If in fact a \mathbf{U} matrix with orthogonal columns is **normalized**, then, furthermore, the sums of squares for the Group effect from each of the $k = 1, \dots, n - 1$ analyses of variance above will **sum** to SS_{GT} , and, similarly, the error sums of squares from each of these will sum to SS_E .

To see this in a special case, consider the Helmert matrix in (3.24) with $n = 4$,

$$\mathbf{U} = \begin{pmatrix} 1 & 0 & 0 \\ -1/3 & 1 & 0 \\ -1/3 & -1/2 & 1 \\ -1/3 & -1/2 & -1 \end{pmatrix}.$$

Each of the columns c_k , $k = 1, \dots, n - 1 = 3$, yields a function of a data vector $\mathbf{c}'_k \mathbf{Y}_{hl}$ that is on a different scale, so that the sums of squares from the individual analyses of variance corresponding to each are not comparable. It is possible to modify each contrast without affecting orthogonality yield a common scale by **normalizing** each column; i.e., divide each column by the square root of the sums of squares of its elements so that the sum of the squares of the modified elements is equal to one. E.g., for \mathbf{c}_1 , the sum of squared elements is $1^2 + (-1/3)^2 + (-1/3)^2 + (-1/3)^2 = 4/3$, yielding the normalized version $\sqrt{3/4}\mathbf{c}_1$; similarly, the normalized versions of \mathbf{c}_2 and \mathbf{c}_3 are $\sqrt{2/3}\mathbf{c}_2$ and $\sqrt{1/2}\mathbf{c}_3$. When all contrasts in an orthogonal transformation are scaled in this way, then they are said to be **orthonormal**. If the orthonormal contrasts are used in the individual analyses above to form the “data,” then the sums of squares from each **do** sum to the overall SS_{GT} and SS_E .

It is **not necessary** to use normalized contrasts to obtain the correct test statistics for each component contrast. The **same** test statistics will result; clearly, although each analysis is on a different scale, the F ratios will be the same, as the normalization factor will **cancel** from numerator and denominator. The orthonormal version of the transformation is often used simply because it leads to the nice, intuitive additive property.

If the component contrasts in the chosen \mathbf{U} matrix are **not orthogonal**, interpretation of the separate tests is more difficult, as the sums of squares are no longer partitioned as above, so that the outcome of one test is related to that of another.

ORTHOGONAL POLYNOMIAL CONTRASTS: As we noted at the outset, the **statistical model** (3.5) does not **acknowledge explicitly** that the response on a given individual likely evolves over **continuous time**. In particular, in the context of the **conceptual framework** in Chapter 2, the model does not incorporate formally an acknowledgment of a **smooth** underlying trajectory, as is apparent in **EXAMPLES 1 – 4** in that chapter.

Accordingly, there is a need to be able to evaluate behavior of the mean response over (continuous) time in the context of the statistical model that acknowledges possible **smooth patterns** of change. For example, in the dental study, we might wish to evaluate whether or not there is a **linear** or in fact **quadratic** trend over time, averaged across genders and whether or not the linear or quadratic trend differs between genders.

This is facilitated by, for n time points, the set of $n - 1$ **orthogonal polynomial contrasts**. These contrasts are based on the premise that, with data at the same n time points, it is possible to fit up to a $(n - 1)$ th degree polynomial in time. Thus, just as the profile and Helmert transformations **decompose** the overall time effect into $n - 1$ contrasts addressing specific differences over time, these contrasts decompose this into **orthogonal components** reflecting the strength of linear, quadratic, cubic, and so on contributions to the saturated $(n - 1)$ th degree polynomial. This is possible for times that are **equally** or **unequally** spaced; we do not present derivations here. For equally-spaced time points, the coefficients of the $n - 1$ orthogonal polynomials are available in many classical statistics texts; for unequally-spaced times, the coefficients depend on the times themselves.

For example, for $n = 4$, there are $n - 1 = 3$ possible (orthogonal) contrasts corresponding to **linear**, **quadratic**, and **cubic** components of the overall smooth trend, which are characterized in the columns of following \mathbf{U} matrix:

$$\mathbf{U} = \begin{pmatrix} -3 & 1 & -1 \\ -1 & -1 & 3 \\ 1 & -1 & -3 \\ 3 & 1 & 1 \end{pmatrix}.$$

It can be verified that the columns of \mathbf{U} are **orthogonal**.

With the appropriate set of orthogonal polynomial contrasts, one can proceed as above to conduct separate hypothesis tests addressing the strength of the linear, quadratic, and so on components of the mean response trajectory over time. The orthogonal polynomial transformation can also be “normalized” as discussed above.

ADJUSTED TESTS: We conclude this section by returning to the assumption embodied in the model (3.5) that the overall aggregate correlation structure is that of **compound symmetry** or at least of **Type H**. As noted previously, if the assumption of Type H does **not** hold, then the usual F tests of the group by time interaction and main effect of time are **invalid** in that they will be **too liberal**.

If the analyst doubts the relevance of this assumption, methods are available to “**adjust**” the usual F tests. We sketch how this approach works without providing technical justification or details.

Define

$$\epsilon = \frac{\text{tr}^2(\mathbf{U}'\mathbf{V}\mathbf{U})}{(n-1)\text{tr}(\mathbf{U}'\mathbf{V}\mathbf{U}\mathbf{U}'\mathbf{V}\mathbf{U})},$$

where \mathbf{U} is any $(n \times n-1)$ matrix whose columns are **normalized orthogonal contrasts**. It can be shown that the constant ϵ defined in this way must satisfy

$$1/(n-1) \leq \epsilon \leq 1$$

and that

$$\epsilon = 1$$

if, and only if, \mathbf{V} is of Type H.

Because the usual F tests are **too liberal** if \mathbf{V} is not of Type H, one suggestion is, rather than compare the F ratios to the usual F critical values with a and b numerator and denominator degrees of freedom, say, compare them instead to F critical values with ϵa and ϵb numerator and denominator degrees of freedom instead. This will make the degrees of freedom **smaller** than usual. It can be verified that, as the numerator and denominator degrees of freedom get **smaller**, the value of the critical value gets **larger**. Thus, the effect of this “adjustment” is to compare F ratios to larger critical values, making it harder to reject the null hypothesis and thus making the test less **liberal**.

- Of course, ϵ is not known, because it depends on the unknown \mathbf{V} matrix. Thus, different adjustments are based on different approaches to **estimating \mathbf{V}** and using the result to estimate ϵ .
- Two such approaches are the **Greenhouse-Geisser** and **Huynh-Feldt** adjustments. Each estimates ϵ in a different way; the Huynh-Feldt estimate is such that the adjustment to the degrees of freedom is not as severe as that of the Greenhouse-Geisser adjustment. These adjustments are based on **asymptotic approximations**, so that it is not necessarily the case that they will lead to valid tests when the numbers of individuals are small

SUMMARY: The spirit of the methods discussed above can be summarized as follows. One adopts a **statistical model** that makes the very specific assumption of **compound symmetry** of the aggregate correlation structure among responses on the same individual.

If this assumption is correct, then familiar analysis of variance methods are available to **test hypotheses** regarding the **pattern of change** over time and mean response averaged across groups. However, the model does not lend itself readily to **estimation** of features of the pattern of change, and the procedures to construct tests to study different features of the pattern are rather **unwieldy**. It is possible to carry out a test of whether or not the compound symmetry assumption is supported by the data; however, the testing procedures are not reliable. Approximate, “adjusted” versions of the tests are available, but these are not necessarily reliable, either.

The bottom line is that a **better approach** might be to start with a more realistic and flexible **statistical model** within which to characterize and evaluate features of the pattern of change. This is the basis for the more **modern** methods we study in later chapters.

3.4 Multivariate repeated measures analysis of variance

We conclude our discussion of classical approaches with a brief overview of **multivariate repeated measures analysis of variance** methods.

MULTIVARIATE MODEL: The set-up and notation are identical to those introduced in Section 3.2; i.e., individuals belong to one of $g \geq 1$ groups, and the response is ascertained on each individual at n time points. There are r_ℓ individuals in each group, indexed by $h = 1 \dots, r_\ell$, for a total of $m = \sum_{\ell=1}^g r_\ell$.

The representation of the **overall population mean response** is the same as for the univariate approach; namely, for the h th individual in the ℓ th group at time j ,

$$E(Y_{h\ell j}) = \mu_{\ell j} = \mu + \tau_\ell + \gamma_j + (\tau\gamma)_{\ell j}.$$

However, as we noted at the beginning of this chapter, these methods are based on a **PA perspective**, so do not acknowledge **among-** and **within-individual** sources of correlation separately and explicitly, and they make **no specific assumption** on the form of the overall pattern of covariance, taking it to be completely **unstructured**.

In particular, the assumed model is

$$Y_{h\ell j} = \underbrace{\mu + \tau_\ell + \gamma_j + (\tau\gamma)_{\ell j}}_{\mu_{\ell j}} + \epsilon_{h\ell j} \quad \text{or} \quad \mathbf{Y}_{h\ell} = \boldsymbol{\mu}_\ell + \boldsymbol{\epsilon}_{h\ell}, \quad \boldsymbol{\epsilon}_{h\ell} \sim \mathcal{N}(\mathbf{0}, \mathbf{V}),$$

where \mathbf{V} is an **arbitrary** covariance matrix with no particular structure; that is, \mathbf{V} is an **unstructured** covariance matrix.

Thus, the model can be summarized as

$$\mathbf{Y}_{h\ell} \sim \mathcal{N}(\boldsymbol{\mu}_\ell, \mathbf{V}), \quad h = 1, \dots, r_\ell, \quad \ell = 1, \dots, q, \quad (3.27)$$

where \mathbf{V} is **completely unstructured**, depending on $n(n+1)/2$ **distinct parameters** (compared to the **two** parameters, σ_b^2 and σ_e^2 , characterizing the compound symmetric structure underlying the univariate methods).

Under model (3.27), it is of course possible to conceive the **same hypotheses** of parallelism, coincidence, and constancy as under the univariate model.

- However, because of the unstructured covariance assumption, it is **no longer possible** to derive straightforward test statistics involving ratios of simple mean squares based on individual response observations.
- Instead, one must view the problem from a **multivariate** perspective and develop testing procedures based on classical **multivariate analysis of variance (MANOVA)** techniques.
- Just as the univariate methods make an assumption, that of **compound symmetry**, for the overall aggregate covariance structure that may be **too restrictive** for many longitudinal data problems, the multivariate methods make an assumption that is usually **too general**.
- Because this assumption is so general, involving $n(n+1)/2$ **covariance parameters**, these procedures are not very **powerful** for detecting departures from null hypotheses of interest.
- Thus, these methods are of **limited practical utility** and are **rarely used** anymore for longitudinal data analysis. Our presentation is meant only to provide an introduction to the basic ideas.

GENERAL MULTIVARIATE PROBLEM: In the longitudinal data setting, the elements of a response vector $\mathbf{Y}_{h\ell}$ are observations on the **same** response over time. In the **general multivariate problem**, the components of $\mathbf{Y}_{h\ell}$ can be, but are **not necessarily**, observations on the same response. Instead, they can be observations on n **different variables**. For example, in a health sciences study, $Y_{h\ell 1}$ might be systolic blood pressure, $Y_{h\ell 2}$ might be diastolic blood pressure, $Y_{h\ell 3}$ might be total cholesterol level, and so on.

In the most general case, then, hypotheses of parallelism or involving **averaging** over the n components of $\boldsymbol{\mu}_\ell$ are **nonsensical**. Instead, the focus is on comparing the means of each of the components **simultaneously** across groups.

That is, the **null hypothesis** of central interest is

$$H_0 : \mu_1 = \cdots = \mu_g \quad (3.28)$$

versus the alternative that at least one of the g population mean vectors differs from the others in at least one component.

We first review the standard approach to testing (3.28) and then discuss how **specialized tests** relevant to the **longitudinal** situation can be developed.

HOTELLING'S T^2 : When $g = 2$, the test statistic for testing H_0 in (3.28) can be viewed as a **generalization** to multivariate response of the usual two-sample t test for scalar response. Here, (3.28) is

$$H_0 : \mu_1 = \mu_2. \quad (3.29)$$

Collecting the sample averages $\bar{Y}_{\cdot\ell j}$ for each component $j = 1, \dots, n$ for group $\ell = 1, \dots, g$ as

$$\bar{\mathbf{Y}}_{\cdot\ell} = \begin{pmatrix} \bar{Y}_{\cdot\ell 1} \\ \vdots \\ \bar{Y}_{\cdot\ell n} \end{pmatrix},$$

the **sample covariance matrix** for group ℓ is

$$\hat{\mathbf{V}}_{\ell} = (r_{\ell} - 1)^{-1} \sum_{h=1}^{r_{\ell}} (\mathbf{Y}_{h\ell} - \bar{\mathbf{Y}}_{\cdot\ell})(\mathbf{Y}_{h\ell} - \bar{\mathbf{Y}}_{\cdot\ell})^T.$$

The sum in this expression is referred to in the multivariate literature as a **sum of squares and cross-products (SS&CP)** matrix. Then the (assumed **common** across groups) **overall covariance matrix \mathbf{V}** is estimated by the **pooled** estimator as in (2.35); with $g = 2$,

$$\hat{\mathbf{V}} = (r_1 + r_2 - 2)^{-1} \{ (r_1 - 1) \hat{\mathbf{V}}_1 + (r_2 - 1) \hat{\mathbf{V}}_2 \}.$$

Analogous to the square of the usual t statistic, the **Hotelling's T^2** statistic is

$$T^2 = (r_1^{-1} + r_2^{-1})^{-1} (\bar{\mathbf{Y}}_{\cdot 1} - \bar{\mathbf{Y}}_{\cdot 2})^T \hat{\mathbf{V}}^{-1} (\bar{\mathbf{Y}}_{\cdot 1} - \bar{\mathbf{Y}}_{\cdot 2}). \quad (3.30)$$

It can be shown that, under model (3.27),

$$\frac{r_1 + r_2 - n - 1}{(r_1 + r_2 - 2)n} T^2 \sim F_{n, r_1 + r_2 - n - 1}.$$

Thus, the test of H_0 may be carried out at level α by comparing this version of T^2 to $\mathcal{F}_{n, r_1 + r_2 - n - 1, \alpha}$. If $n = 1$, the test reduces to the usual two-sample t test.

As an example, consider the dental study, for which $r_1 = 11$ (girls), $r_2 = 16$ (boys), $n = 4$, and

$$\bar{\mathbf{Y}}_{..1} = (21.182, 22.227, 23.091, 24.091)^T,$$

$$\bar{\mathbf{Y}}_{..2} = (22.875, 23.813, 25.719, 27.469)^T.$$

Using the estimated sample covariance matrices for each group in (2.33) and (2.34) and the resulting pooled estimate, it is straightforward to obtain

$$\frac{r_1 + r_2 - n - 1}{(r_1 + r_2 - 2)n} T^2 = 3.63,$$

which under (3.27) has an F distribution with 4 and 22 degrees of freedom; $\mathcal{F}_{4,22,0.05} = 2.816$, leading to rejection of H_0 at level $\alpha = 0.05$.

Of course, as noted in Section 2.6, the assumption of a common overall pattern of covariance for boys and girls, embodied in this procedure, does not seem to be supported by the data. Moreover, the data support a pattern for each group that, although different for each group, is approximately **compound symmetric**, suggesting that a more powerful test could be developed.

Regardless, this hypothesis test does not address the questions of interest here. Although the result suggests there is evidence that the overall population means differ between genders, this test offers no insight into **how** nor into how the **pattern of change** differs, so is relatively **useless**.

As discussed earlier, defining as in (3.17)

$$\mathcal{M} = \begin{pmatrix} \mu_{11} & \cdots & \mu_{1n} \\ \mu_{21} & \cdots & \mu_{2n} \end{pmatrix}$$

and letting $\mathbf{C} = (1, -1)$ and $\mathbf{U} = \mathbf{I}_n$, H_0 in (3.29) can be expressed as $H_0 : \mathbf{CMU} = \mathbf{0}$. We will return to this representation shortly.

ONE-WAY MANOVA: When $g > 2$, a **multivariate version** of the usual analysis of variance for a **one-way layout** can be constructed and test statistics derived for testing (3.28) as follows. The usual analysis of variance where the response is scalar involves sums of squares and mean squares for “Among Groups” and “Among-Unit Error,” and the ratio of the latter yields the usual test statistic. Here, with **multivariate** response these are replaced by analogous **SS&CP matrices** as follows.

Again let $\bar{\mathbf{Y}}_{..j}$ be the sample average of all observations across all individuals and groups of the j th component of \mathbf{Y}_{hej} , and define the **overall sample mean vector**

$$\bar{\mathbf{Y}}_{..} = \begin{pmatrix} \bar{\mathbf{Y}}_{..1} \\ \vdots \\ \bar{\mathbf{Y}}_{..n} \end{pmatrix}.$$

Then construct the **MANOVA table** as

Source	SS&CP	DF
Among Groups	$\mathbf{Q}_H = \sum_{\ell=1}^g r_{\ell}(\bar{\mathbf{Y}}_{\cdot\ell} - \bar{\mathbf{Y}}_{\cdot\cdot})(\bar{\mathbf{Y}}_{\cdot\ell} - \bar{\mathbf{Y}}_{\cdot\cdot})^T$	$g - 1$
Among-unit Error	$\mathbf{Q}_E = \sum_{\ell=1}^g \sum_{h=1}^{r_{\ell}} (\mathbf{Y}_{h\ell} - \bar{\mathbf{Y}}_{\cdot\ell})(\mathbf{Y}_{h\ell} - \bar{\mathbf{Y}}_{\cdot\ell})^T$	$m - g$
Total	$\mathbf{Q}_H + \mathbf{Q}_E = \sum_{\ell=1}^g \sum_{h=1}^{r_{\ell}} (\mathbf{Y}_{h\ell} - \bar{\mathbf{Y}}_{\cdot\cdot})(\mathbf{Y}_{h\ell} - \bar{\mathbf{Y}}_{\cdot\cdot})^T$	$m - 1$

It can be verified that

$$\mathbf{Q}_E = (r_1 - 1)\hat{\mathbf{V}}_1 + \cdots + (r_g - 1)\hat{\mathbf{V}}_g,$$

so that

$$\hat{\mathbf{V}} = \mathbf{Q}_E / (m - g).$$

Because the entries in the MANOVA table are **matrices**, it is not straightforward to construct a **unique generalization** of the usual analysis of variance F ratio that can be used to test H_0 in (3.28). Clearly, one would like to compare the “**magnitudes**” of the SS&CP matrices \mathbf{Q}_H and \mathbf{Q}_E , but there is no one way to do this. Several statistics have been proposed.

- The most commonly discussed statistic is **Wilks’ lambda**, which can be motivated **informally** as follows. Letting SS_G and SS_E be the usual analysis of variance Among-Groups and Among-Unit Error sums of squares, the familiar F ratio is

$$\frac{SS_G / (g - 1)}{SS_E / (m - g)}.$$

Thus, in the scalar case, H_0 is rejected when SS_G / SS_E is “large.” This is equivalent to rejecting for large values of $1 + SS_G / SS_E$ or small values of

$$\frac{1}{1 + SS_G / SS_E} = \frac{SS_E}{SS_G + SS_E}.$$

For the multivariate problem, the Wilks’ lambda statistic is the **analog** of this quantity,

$$T_W = \frac{|\mathbf{Q}_E|}{|\mathbf{Q}_H + \mathbf{Q}_E|}.$$

One rejects H_0 for “small” values of T_W .

- The **Lawley-Hotelling trace** rejects H_0 for large values of

$$T_{LH} = \text{tr}(\mathbf{Q}_H \mathbf{Q}_E^{-1}).$$

- Other statistics are **Pillai’s trace** and **Roy’s greatest root**, which we do not present.
- None of these approaches has been shown to be superior to the others in general. All are **equivalent** to using the Hotelling T^2 statistic in the case $g = 2$.

A full discussion of these methods is beyond our scope. For general g and n , the sampling distributions of (functions of) these test statistics may or may not be derived **exactly**. Thus, except in certain special cases where this is possible (see Johnson and Wichern, 2002), the sampling distributions are **approximated** by the F or other distributions.

As in the case of $g = 2$, the hypothesis (3.28) can be expressed in the form $H_0 : \mathbf{C}\mathbf{M}\mathbf{U} = \mathbf{0}$ for appropriate choice of \mathbf{C} and $\mathbf{U} = \mathbf{I}_n$. For example, for $g = 3$,

$$\mathbf{M} = \begin{pmatrix} \mu_{11} & \cdots & \mu_{1n} \\ \mu_{21} & \cdots & \mu_{2n} \\ \mu_{31} & \cdots & \mu_{3n} \end{pmatrix}, \quad \mathbf{C} = \begin{pmatrix} 1 & -1 & 0 \\ 1 & 0 & -1 \end{pmatrix}.$$

Again, as for $g = 2$, in the **longitudinal data** situation, testing (3.28) does not really address the questions of scientific interest, which usually focus on the **pattern of change**. We now consider how tests of **parallelism** (group by time interaction) and **constancy** (main effect of time) are developed under model (3.27), so under the assumption of an **unstructured covariance matrix** common across groups.

PROFILE ANALYSIS: In the context of **repeated measurement data** where the n components of a data vector are repeated observations on the same response, conducting appropriate such **multivariate tests** for parallelism and constancy is referred to as **profile analysis**.

Consider first the hypothesis of **parallelism** or **group by time interaction**. As in the univariate case, under the usual constraints (3.11), this hypothesis is

$$H_0 : \text{all } (\tau\gamma)_{\ell j} = 0.$$

For $g = 2$ and $n = 3$, we expressed this hypothesis in the form

$$H_0 : \mathbf{C}\mathbf{M}\mathbf{U} = \mathbf{0},$$

with \mathbf{C} and \mathbf{U} given in (3.19) and (3.20). For general g and n , with \mathbf{M} as in (3.17), it follows that, if $\mathbf{1}_p$ denotes a column vector of 1s of length p , then choosing

$$\mathbf{C} = \begin{pmatrix} \mathbf{1}_{g-1} & -\mathbf{I}_{g-1} \end{pmatrix} \quad (g-1 \times g), \quad \mathbf{U} = \begin{pmatrix} \mathbf{1}_{n-1}^T \\ -\mathbf{I}_{n-1} \end{pmatrix} \quad (n \times n-1) \quad (3.31)$$

yields the null hypothesis of parallelism.

We now explain informally the reason for writing hypotheses in the form $H_0 : \mathbf{C}\mathbf{M}\mathbf{U} = \mathbf{0}$, which provides the basis for deriving test statistics for the hypotheses of parallelism and constancy. Consider again the hypothesis (3.28), $H_0 : \mu_1 = \cdots = \mu_g$, which can be written for general g and n in this form, with \mathbf{C} as in (3.31) and $\mathbf{U} = \mathbf{I}_n$. It can be shown that the SS&CP matrices \mathbf{Q}_H and \mathbf{Q}_E on which the various test statistics are based can be expressed in a general form in terms of \mathbf{C} , \mathbf{M} , and \mathbf{U} , as follows.

Let \mathbf{A} be the $(m \times q)$ matrix whose **rows** each correspond to one of the m individuals in the data set as follows. For individual (h, ℓ) , the corresponding row is a $(1 \times q)$ vector of all 0s except for a 1 in the ℓ th position; for example, with $g = 3$, for an individual in group 2, the corresponding row is

$$\mathbf{a}_2 = (0, 1, 0),$$

say. Premultiplying the $(g \times n) = (3 \times 4)$ matrix

$$\mathbf{M} = \begin{pmatrix} \mu_{11} & \mu_{12} & \mu_{13} & \mu_{14} \\ \mu_{21} & \mu_{22} & \mu_{23} & \mu_{24} \\ \mu_{31} & \mu_{32} & \mu_{33} & \mu_{34} \end{pmatrix}$$

then yields

$$\mathbf{a}_2 \mathbf{M} = (\mu_{21}, \mu_{22}, \mu_{23}, \mu_{24}) = \boldsymbol{\mu}_\ell^T.$$

That is, the matrix \mathbf{A} “picks off” the mean vector (row of \mathbf{M}) corresponding to the group to which each individual belongs; a little thought reveals that, in general, \mathbf{A} will have r_1 rows $\mathbf{a}_1 = (1, 0, \dots, 0)$, r_2 rows $\mathbf{a}_2 = (0, 1, 0, \dots, 0)$, \dots , r_g rows $\mathbf{a}_g = (0, 0, \dots, 1)$. It can be deduced that the model $\mathbf{Y}_{h\ell} = \boldsymbol{\mu}_\ell + \epsilon_{h\ell}$ can be written succinctly as

$$\mathbf{Y} = \mathbf{A}\mathbf{M} + \boldsymbol{\epsilon}, \quad (3.32)$$

where \mathbf{Y} is the $(m \times n)$ matrix with rows $\mathbf{Y}_{h\ell}^T$, and similarly for $\boldsymbol{\epsilon}$.

Then it is an exercise in matrix algebra to show that, with \mathbf{C} as in (3.31) and $\mathbf{U} = \mathbf{I}_n$, the SS&CP matrices \mathbf{Q}_H and \mathbf{Q}_E in the MANOVA table can be written as

$$\mathbf{Q}_H = (\mathbf{C}\widehat{\mathbf{M}}\mathbf{U})' \{ \mathbf{C}(\mathbf{A}'\mathbf{A})^{-1} \mathbf{C}' \}^{-1} (\mathbf{C}\widehat{\mathbf{M}}\mathbf{U}), \quad \mathbf{Q}_E = \mathbf{U}' \mathbf{Y}' \{ \mathbf{I}_n - \mathbf{A}(\mathbf{A}'\mathbf{A})^{-1} \mathbf{A}' \} \mathbf{Y} \mathbf{U}, \quad (3.33)$$

where $\widehat{\mathbf{M}} = (\mathbf{A}'\mathbf{A})^{-1} \mathbf{A}'\mathbf{Y}$.

A technical justification of (3.33) can be found in Vonesh and Chinchilli (1997, p. 50), who show that this representation and the form of the Wilks' lambda statistic T_W can be derived via **maximum likelihood** under model (3.32) and the normality assumption (3.27).

It can be shown that test statistics for **other hypotheses**, such as parallelism and constancy, can be derived by substituting the relevant \mathbf{C} and \mathbf{U} , such as those in (3.31), into (3.33) to yield versions of the SS&CP matrices \mathbf{Q}_H and \mathbf{Q}_E that can be used to construct versions of any of test statistics above, such as Wilks' lambda T_W , that address the corresponding hypotheses. Depending on g and n , these tests may be exact or approximate.

Classically, **profile analysis** has been carried out in practice as follows:

- The test of primary interest in longitudinal settings, that of **parallelism** or **group by time interaction**, is carried out first by taking \mathbf{C} and \mathbf{M} to be as in (3.31) and constructing the desired test statistic.
- If the hypothesis of parallelism is **not rejected**, the test of **coincidence** is carried out; this is the usual MANOVA test described above, with \mathbf{C} as in (3.31) and $\mathbf{U} = \mathbf{I}_n$. If the mean response profiles are **not parallel**, then this test seems difficult to justify. If the profiles **are parallel**, then this test can be **refined**. With the additional assumption of parallelism, this is equivalent to a test with \mathbf{C} as in (3.31) and $\mathbf{U} = \mathbf{1}_n/n$, which leads to exactly the hypothesis corresponding to the **main effect of group** in the univariate analyses discussed in Section 3.2. In fact, it can be shown that any of the test statistics discussed above constructed using this \mathbf{C} and \mathbf{U} reduces to the F statistic for the main effect of group in the univariate analysis.
- If the hypothesis of parallelism is **not rejected**, the test of **constancy** of the mean profiles over time is carried out. It can be shown that this test corresponds to taking \mathbf{U} as in (3.31) and $\mathbf{C} = \mathbf{I}_g$. As with the test for coincidence, if the profiles are **not parallel**, then testing whether they are **constant** over time seems inappropriate. Under the additional assumption of parallelism, a refined test can be constructed corresponding to \mathbf{U} as in (3.31) and $\mathbf{C} = \mathbf{1}_g^T/g$. Unlike for the refined test of coincidence, the resulting multivariate tests are **different** from each other and from the univariate test in Section 3.2.

SUMMARY: Our discussion of multivariate analysis of variance methods has been **deliberately brief** because, as hopefully is evident from their formulation, these methods are clearly of **limited utility** in longitudinal data analysis. The focus on **hypothesis testing** and the assumption of a **common but unstructured** overall covariance matrix are key drawbacks. It is important, however, to have a basic understanding of these methods so that the appeal of the more modern methods we discuss henceforth can be fully appreciated.